

STAT 301

A/B Testing (2-sample Hypothesis Tests)

- test to compare the param of two pop: control (A) & variation (B)
 - basically 2-sample hypothesis test (t -test if population sd σ is unknown, z -test if known)
 - need to randomly assign ppl to make it an experiment
- hypothesis test: always hypothesizing about the population param
 - H_0 is the null hypothesis, it refers to the status quo (i.e there's no change in engagement) - **it's always an equality**
 - H_1 is the alternative hypothesis, it's the researcher's hypothesis of interest (i.e there's an increase in engagement)

Means

- code to test $H_a : \mu_1 > \mu_2$ (mean of sample 1 vs sample 2)

```
# this is testing is x > y (so x - y > 0)
result <- tidy(t.test(x = sample1, y = sample2,
  alternative = "greater", # or "two.sided" or "lesser"
  var.equal = FALSE)) # TRUE if assumed equal var
```

- from result – we can get the p -val, if it's less than significance level, we reject (p is low, null must go – it's evidence against the null)
- another way to get the p -value is **bootstrapping**, the code is

```
obs_test_stat <- tiktok_sample %>% # test stat from og sample
  specify(formula = response_var ~ explanatory_var) %>%
  calculate(stat = "diff in means", # it's new - current
    order = c("new", "current")) %>% pull()
pvalue = tiktok_sample %>% # get p-value from test stats
  specify(formula = response_var ~ explanatory_var) %>%
  hypothesize(null = "independence") %>%
  generate(reps = 1000, type = "permute") %>%
  calculate(stat = "diff in means",
    order = c("new", "current")) %>% # compare resample
  get_p_value(obs_stat = obs_test_stat, # test stat against
    direction = "greater") %>% pull() # og test stat
```

- errors in hypothesis testing
 - Type I error**: reject H_0 when H_0 is true, denoted α
 - Type II error**: fail to reject H_0 when H_0 is false, denoted β
 - α = significance level is set by experiment designer (i.e is **fixed**)
- power**: prob of rejecting H_0 when H_0 is False \rightarrow $\text{Power} = 1 - \beta$
 - larger (sample size, effect size, α) = bigger power
 - effect size is the difference between hypothesized and true mean
 - smaller variance = bigger power

Proportions

- when doing A/B testing for proportions - it's two-sample z -test
 - ex. $H_0 : p_A = p_B$ vs. $H_1 : p_A > p_B$
- if the samples are dependent - do pairwise test

```
pairwise_comparisons <- pairwise.prop.test(
  x = successes, # x is vector of counts of successes
  n = trials,    # n is vector of counts of trials
  p.adjust.method = "bonferroni", alternative = "two.sided")
```

- if just doing two-sample proportion test

```
result <- prop.test(x = c(success_group1, success_group2),
  n = c(total_group1, total_group2),
  alternative = "two.sided")
```

- Bonferroni adjustment: needed for pairwise to prevent increased Type I error (see later)
 - how: do $\alpha_{\text{new}} = \alpha / \#$ of comparisons (ex. 0.05/10)

Early Stopping & Principled Peeking

- some platforms allow users to continuously monitor the p -values and CI as you get more people (sample size changes/increases)
- big point: stopping an experiment & rejecting H_0 as soon as the p -value falls below specified α can drastically inflate Type I error
- principled peeking: methods to solve the early-stopping problem
 - a basic way is Bonferroni method – see above
 - other is Pocock & O'Brien-Fleming, but all of them can be thought of as reducing α or increasing critical val to reject less

- common experiment to demonstrate FWER error is A/A testing
 - you have 2 samples and you **know** their parameters are identical, so every time you reject, you know you're incorrectly rejecting
- code for principled peeking (A/A testing)

```
crit_unadj <- qt(1 - alpha, n) # getting critical value
crit_bonferroni <- qt(1 - (alpha/num_experiments), n)
exp_design <- gsDesign(
  k = num_experiments, # Number of interim analyses
  test.type = 1,        # One-sided test
  alpha = 0.05,         # Significance level
  beta = 0.2,           # 1 - power
  sfu = "Pocock"/"OF",  # Pocock or O'Brien Fleming
  n.fix = 1              # This adjusts the sample size
  # calculation; for A/A testing, effect size = 0
)
crit_adj_pocock_or_of <- exp_design$upper$bound
```

- only O'Brien changes crit_val as n inc (starts very high)
- Pocock is less conservative than Bonferroni
- α : $\text{unadj} > \text{Pocock} > \text{Bonferroni}$ (opposite for critical-value)
- all of these affect (decreases) power of the test ($\alpha \propto 1/\text{power}$)
- Bonferroni bad for seq test bc too conservative & assume indep

Simple Linear Regression

Intro

- simple here means only 1 input variable
 - Y: response variable, dependent variable, output
 - X: explanatory variables, independent variables, features
- the model:

$$Y = \beta_0 + \beta_1 + \varepsilon \quad E[\varepsilon | X] = 0$$

- ε is the error term - contains all other factors affecting Y
- β_0 is the (true) intercept variable, β_1 is the (true) slope variable
 - we will estimate these using `lm()` and denote it $\hat{\beta}_0, \hat{\beta}_1$

Estimation of LR Line

- methods we use to get the linear regression line is **Least Squares** (LSE) method – minimize sum of squares of the residuals (ε_i^2)
- code: do `lm(response_var ~ explanatory_var, data)`
 - we get back a model that has the estimated values $\hat{\beta}_0$ and $\hat{\beta}_1$
- plotting the line: add `geom_smooth(method="lm", se=FALSE)`
- say “for every 1 unit increase of x , y is expected to inc/dec by β_1 ”

SLR Inference (Hypothesis Test + CI)

- variation: estimates of β_0 and β_1 depends on our sample (they may vary between our samples)
 - variation of these estimates from sample to sample is called standard error (note: we can't take multiple samples IRL)
 - in reality: take 1 sample & compute SE via sampling dist
- hypothesis testing: question is “is the input variable linearly associated with the response”

$$H_0 : \beta_i = 0$$

$$H_A : \beta_i \neq 0$$

- ex. null is saying the slope β_1 is 0, corresponds to the assumption that there is no linear relationship between X and Y
- we can check this using `tidy(lm_model, conf.int=TRUE)`
 - give 95% CI as well (change `conf.level` as needed)
 - this is t -test: if p -value is low, reject the null
 - if we reject, we say the variables are statistically associated
- note: high p -value means we don't have strong evidence to reject $\beta_i = 0$, doesn't mean that $\beta_i = 0$ or that they're not associated
- sampling distribution: there are 2 ways to get it
 - theoretical: this is what `lm()` does (**and what we focus on**)

$$T = \frac{\hat{\beta}_1 - 0}{\text{SE}(\hat{\beta}_1)} \sim t_{n-k} \quad k = \# \text{ of regression param}$$

- use `qt()` like above to get test stat and p -value

- bootstrapping: resample w/ rep from OG sample, calc estimate, get list of estimates, get our estimated sampling dist

```
bootstrap_sample <- replicate(num_replicate, {
  sample_n(og_sample, sample_size, replace=TRUE) %>%
  lm(y_var ~ x_var, data = .) %>% . $coef
})
bootstrap_sample <- data.frame(
  boot_intercept = bootstrap_sample[1, ],
  boot_slope = bootstrap_sample[2, ]
)
# or instead you can use infer
bootstrap_slope_infer <- original_sample %>%
  specify(formula = y_var ~ x_var) %>%
  generate(reps = 1000, type = "bootstrap") %>%
  calculate(stat = "slope")
head(bootstrap_slope_infer)
visualize(bootstrap_slope_infer, bin=30) # histogram
```
- larger `num_replicate` means better approximation of dist
- larger og sample size means sampling dist tighter & more smooth
- note: the sampling dist is about the statistics (i.e stuff with hats) and SLR sampling dist is related to $\hat{\beta}_i$
- confidence interval
 - theoretical approach: $CI = \hat{b} \pm SE(\hat{b}) \times t_{\alpha/2, n-k}$ (use `tidy()`)
 - bootstrap (percentile) approach: get the sampling distribution, run `quantile()` or `infer` on it to get the 95% CI

```
# using quantile
CI <- bootstrap_sample %>%
  summarize(lower = quantile(boot_slope, 0.025),
    upper = quantile(boot_slope, 0.975))

# using infer
percentile_ci <- bootstrap_slope_infer %>%
  get_confidence_interval(type = "percentile", level=0.95)
```
 - interpretation: given a CI, we say “we are 95% **confident** that the true coefficient is in the given range”
 - note: CI not including 0 also show statistical significance

Multiple Linear Regression

Additive models (Continuous Variable)

- we use the + in `lm()` function to indicate additive models

```
# povertyPercent and PctPrivateCoverage are continuous
MLR_poverty_coverage <- lm(formula=TARGET_deathRate ~
  # povertyPercent + PctPrivateCoverage, US_cancer_data)
```

- after calling `tidy()` on these we'll get p -value and estimates of each and we can interpret them **separately**
 - i.e we can reject the null btw y and x_1 or reject the null btw y and x_2 but nothing about their combination
 - ex. if estimate for $x_1 = 0.5$, means that for one-unit inc in x_1 , predict a 0.5 unit inc in y , holding all other variables constant

Additive Models (Categorical + Continuous Variables)

- here, we'll look at adding 1 cont. variable and 1 categorical variable
- bc model is additive (no interaction), **they share a common slope**
- to rep categorical var - we'll use dummy variables (turn on or off)
 - if a categorical var has k levels, you need $k - 1$ dummy var
- case study: continue with the example of the 2 states

$$Y_i = \beta_0 + \beta_1 \text{isWashington} + \beta_2 \text{povertyPercent} + \varepsilon_i$$

here, Indiana is called the **baseline** (it's alphabetical)

$$\text{in Indiana : } Y_i = \beta_0 + \beta_1(0) + \beta_2 \text{povertyPercent} + \varepsilon_i \\ = \beta_0 + \beta_2 \text{povertyPercent} + \varepsilon_i$$

$$\text{in Washington : } Y_i = \beta_0 + \beta_1(1) + \beta_2 \text{povertyPercent} + \varepsilon_i \\ = (\beta_0 + \beta_1) + \beta_2 \text{povertyPercent} + \varepsilon_i$$

- code: `lm(y ~ state + povertyPercent, data = cancer_data)`
- interpretation
 - β_0 is intercept of the **reference** line.
 - β_1 (coefficient of the dummy variable) is the **difference** between intercepts of both lines (inc/dec when going from Indy to Wash)
 - β_2 is the **common** slope of both lines

Interactive Model (Categorical + Continuous Variable)

- say we believe slope changes btw states too (add interaction terms)
- can interpret following as 2 separate LR in 1 eq, representing 2 lines

$$Y_i = \beta_0 + \beta_1(\text{isWashington}) + \beta_2(\text{povertyPercent}) + \beta_3(\text{isWashington} \times \text{povertyPercent}) + \varepsilon_i$$

- in Indiana: `isWashington = 0`

$$Y_i = \beta_0 + \beta_1(0) + \beta_2(\text{povertyPercent}) + \beta_3(0) + \varepsilon_i = \beta_0 + \beta_2(\text{povertyPercent}) + \varepsilon_i$$

- in Washington

$$Y_i = (\beta_0 + \beta_1) + (\beta_2 + \beta_3)\text{povertyPercent} + \varepsilon_i$$

- code: `lm(y ~ state * povertyPercent, data = cancer_data)`
- interpretation
 - β_0/β_2 is the intercept/slope of the **reference** line
 - β_1/β_3 is the difference in intercept/slopes between both lines (ex. for Wash the intercept is $(\beta_0 + \beta_1)$ and the slope is $(\beta_2 + \beta_3)$)
- note: choice of baseline doesn't affect \hat{y} , it changes interpretation (ex. can only speak about stat significance of slope for the baseline)
- largest model: number of scenarios \times number of continuous vars

LR: Model Assumption & Diagnostic

- Linear relation: is it LR
 - "linear"= linear in terms of linear combination of the variables
 - consequence: regression line might not fit (represent) data well
 - diagnostic: look at residuals vs. fitted values plot - should look random without a pattern for a good fit
- Errors are independent
 - can be assessed from the design of the study \rightarrow randomization and repeated experiments ensure the independence of errors
 - consequence: estimates may be underestimated, which can lead to falsely significant p-values
 - ex. time series data is not independent
- Is the conditional distribution of the errors Normal?
 - note: errors don't have to be Normal for valid inference result \rightarrow approximate sampling dist via CLT (large n) or bootstrapping
 - diagnose: use Q-Q plot - straight line on QQ plot is good
 - consequence: the estimates of params are not heavily affected, but the SE are large (which we rely on to create distributions)
 - fix: variables transformations can be used as possible "remedies"
- Equal variance of the error terms (aka homoscedasticity)
 - diagnose: use residuals-fitted plot - funnel shape = bad
 - consequence: the estimates will be similar to the true population parameters but their estimated standard errors will be inflated
 - fix: try transforming the variable
- Multicollinearity: occurs when (some of) input var are correlated
 - when that happens, the information of one variable can be masked by another variable carrying correlated data
 - diagnose: check correlation (`cor(x,y)`) bt var via pairwise plots
 - diagnose: detect multi-col using variance inflation factor (VIF)
 - if $VIF_j \gg 1$, there's multicollinearity w/ x_j in the data
 - code: `vif(lr_model)` – remove the ones w the highest score
 - consequences:
 - inflates SE of the regression estimator
 - cause instability in the regression coefficients (small changes to data result in large changes in the coefficients)
 - mask true importance of predictors by inflating coefficients/significance of correlated var (unreliable results)
 - fix: select which variable (among the correlated ones) to keep

Statistical Design and Causality

- confounding factors: **confounder** is a variable that causes changes in both the response **and** at least one input variable
 - consequences: distort observed relationship btw vars, lead to false conclusions (inaccurate estimates, doesn't generalize well)
 - fix: different approach bt observational & experiment (see later)

- confounder example: explanatory variable is exercise amount, response variable is risk of heart disease, confounding variable is diet
 - healthier diets correlate with more exercise and lower heart disease risk; but we may incorrectly attribute heart disease risk reduction to exercise alone; accounting for diet in analysis clarifies exercise's true effect on heart disease
- casual inference: establishing causal effects (challenging)
 - depends on how data is collected & stat method used for analysis
- Completely Randomized Design (CRD): experimental units are randomized throughout the data layout \rightarrow can establish causal rela
 - observed and unobserved confounders are balanced, on average
 - considered the gold standard design for causal inference
- Randomized Block Design (RBD): splits experimental units into homogeneous blocks to remove variation from nuisance factors, then randomly assigns treatments to each block
 - ex. subjects of similar age groups are blocks, grouping by sex
 - in RBD, only observed confounders are balanced so only average treatment effects can be estimated (via appropriate methods)
- observational data: we collect data by measuring variables or surveying members without applying any treatment to them
 - treatments not controlled by design - **causal effects can not be naively established** (need to include confounder in model)

Goodness of Fit & Nested Models

- the residual: defined as $\text{res}_i = y_i - \hat{y}_i$
 - where \hat{y}_i is the predicted value $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$
 - in R we can do `lm.fitted` and `lm.res`
- quantiles used
 - ESS (explained SS): if model better than nothing, this would be large, measures how much var in data explained by LR
 - RSS (residuals SS): our LR algo try to minimize this
 - TSS (total SS): res SS from the null model, $TSS = ESS + RSS$
- goodness of fit is asking "is our model better than the null model"
 - null model is the intercept only model (where the intercept is estimated by the sample mean \bar{y} bc $E[Y] = \bar{y}$)
 - null model in R is `lm(response ~ 1)`
- coefficient of determination R^2** : interpreted as the proportion of variance of the response explained by the model
 - math: $R^2 = 1 - RSS/TSS = ESS/TSS$
 - ranges from 0 to 1, closer to 1 is better
 - in R, can get all summary statistics by doing `glance(model)`
 - limitation: can't use to say how good model is at extrapolating
 - limitation: inc as new var added to the model (regardless of relevance) - can't be used to compare models of diff size
- adjusted R^2** : accounts for model complexity (penalize added var)
 - math: adjusted $R^2 = 1 - \frac{RSS/(n-p)}{TSS/(n-1)}$
 - p is number of coefficients (variables) and n is sample size
 - like R^2 , closer to 1 is better
- Residual Standard Error RSE**
 - math: $RSE = \sqrt{\frac{1}{n-p} \times RSS}$
 - called **sigma** in `glance(model)`
 - estimates the standard deviation of the error term ε
 - smaller is better
- Mean Squared Error MSE**
 - math: $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
 - often used in cases where the focus is on prediction accuracy
 - training MSE: `.resid` column in the `augment(model)` output
 - testing MSE: compute on new data and predicted value, use to evaluate out-of-sample prediction performance
 - $RMSE = \text{root mean square error} = \sqrt{MSE}$
- about `glance()`: the H_0 it's testing is if **all** predictors are 0 (ANOVA - i.e our model is not better than the null)

```
my_glance <- data %>% summarize(
  RSS = sum(lm_model$resid^2),
  TSS = sd(data$response)^2 * (n - 1),
  R2 = 1 - RSS/TSS,
  adjR2 = 1 - (RSS/(n-2)) / (TSS/(n-1)),
  sigma = sqrt(RSS/(n-2)))
```

The F-Test (ANOVA)

- used to compare model of different size (i.e full vs reduced model)
- hypothesis test: testing if any of the additional variables are zero

$$H_0 : \beta_{q+1} = \beta_{q+2} = \dots = \beta_s = 0$$

$$H_1 : \text{at least one } \beta_j \neq 0 \quad (\text{for } j = q + 1, q + 2, \dots, s)$$

- in R (NB: `glance()` is also doing F-test but against the null)

```
lm_red <- lm(protein ~ 1, dat_3genes) # null
lm_full <- lm(protein ~ gene + mrna, dat_3genes) # full
anova(lm_red,lm_full)
```

 - note: ANOVA can compare b/t any models, not just the null
 - small p -value = stat evidence that complex model provides better fit than the simpler model (explains more variability)
- t vs F test: t to assess the sig of individual coeff; F to assess the overall sig of model or compare nested models (full vs reduced)
 - note: equivalent when only 1 coeff is added btw full and reduced
 - ex. if comparing $y = \beta_0 + \beta_1$ vs $y = \beta_0$, then `tidy(full_model)`, `glance(full_model)`, `anova(full_model)` are all the same
- from middy: if asked to choose b/t `anova()`/`tidy()`, choose `anova()` if categorical var has > 2 levels, `tidy()` otherwise

Variable Model Selection

Generative Model

- may want to identify the most relevant var to build model
 - need to choose an evaluation metric, this depends on the goal
- forward selection: goal is to select the best model from among all possible models of varying sizes (`num_var = 2number of var`)
 - iteratively add var to the model starting from an intercept-only
 - proceeds to evaluate models of increasing size (start w models w 1 variable, then 2, and so on, until reaching the full model)
 - at each step/size, best model picked based on a metric like RSS
 - end: desired size reached or no improvement after add new var
- popular metrics for model selection: adjusted R^2 , test mean squared error (MSE), C_p (AIC), or Bayesian information criterion (BIC)
 - C_p , AIC , BIC : smaller = better
- note: these selection algorithms are greedy and might not give globally best model (also might vary between runs)
- R code: for forward selection you can just do `method="forward"`

```
backward_sel <- regsubsets( # NOTE: code for both gen/pred
  x = response_var ~ explanatory_var,
  nvmax = max_size_of_model,
  data = test_set, # or training if generative
  method = "backwards" # or "forward", or "exhaustive"
)
reg_summary <- summary(backwards_sel) # lowest Cp = best model
cp_min = which.min(reg_summary$cp) # get model with min CP
selected_var <- names(coef(backwards_sel, cp_min))[-1] # names
```

Predictive Model

- models meant for inference is called generative models, models to predict is called predictive models – use different metrics
 - generative models: use R^2 , adjusted R^2 , MSE , F -test, etc
 - predictive model: use MSE, RMSE, R^2
 - for all these, can get metrics on the test set if you have one
 - can also do C_p , AIC, BIC on test set
 - note: can still use these metrics for predictive w/o test set
- test MSE: predict on test set, then calculate residuals/MSE
- test R^2 : do $R^2 = \text{corr}(\hat{y}_{\text{test}}, y_{\text{test}})$ on the test set
 - interpretation: it's measure of *corr* between true outcomes & model's predictions, no longer how much var model explains
- code: to get metrics on test set

```
test_metrics <- metrics(data, truth = target_col_name,
  estimate = prediction_col_name) # attach preds onto og DS
```

Uncertainty of Predictions

- predictions are random variables and thus have uncertainty
 - note: both CIP and PI are centered around the **same value**
- confidence intervals for prediction (CIP): range for the expected (average) val of Y for a given value of X (centered at fitted val \hat{Y}_i)
 - uses the regression line which introduces uncertainty (i.e $\hat{\beta}_i$)
 - code: `predict(lm_model, newdata, interval="confidence", level = 0.95, se.fit=TRUE)`
 - interpretation: with 95% confidence, the **expected** (average) value of data point with X fall between this range
- prediction interval (PI): range within which we can expect to find the actual observed value (like from og data) of Y for a given X
 - same uncertainty as the CIP plus uncertainty of ε_i so while it's also centered at \hat{Y}_i but is **wider than CIP**
 - code: `predict(lm_model, newdata, interval="prediction")`
 - interpretation: with 95% conf, the Y val for data point with X val will be in this range (note: don't say expected anymore)

Predictive Modelling with LR

- split data into training & testing set (imagine data had ID col)

```
training_data <- sample_n(data,size=nrow(data)*0.7,replace=F)
training_data2 <- slice_sample(data, prop = 0.7) # another way
testing_data <- anti_join(data, training_data, by = "ID")
```
- build model using training data then run model on testing data

```
model <- lm(..., data = training_data)
predictions <- predict(model, newdata = testing_data)
RMSE <- rmse(preds = predictions, # or can use metrics()
             actuals = testing_data$target) # get test RMSE of model
```

Inference After Model Selection

- if you use the same dataset to model select **and** make inference, you risk overfitting \rightarrow biased estimates and inflated Type I err
 - Type I mean saying model not better than null (when it is)
 - you need to split - training for model selection (additionally using CV), then testing set to see performance
- usually there's trade-off between performance and interpretability
- overfitting: means model won't generalize well to unseen data

Regularized Regression Methods

- by using regularization: while training, they'll assign certain coeff weights of 0 (effectively removing - does both training & selecting)
- regularized models like Ridge and Lasso adds a penalty to loss function which shrink the coefficients towards 0
 - Ridge (L2 penalty) doesn't set coeffs to 0, but removes highly correlated variables (Lasso also works on high corr ds)
 - Lasso (L1 penalty) will set coeffs to 0, can use for model select
- regularization biases the coeffs by imposing penalty on their magnitude (trade-off: increase bias to decrease variance)
 - this means sampling dist doesn't center on true parameter val
- standardizing input var is necessary b/c methods sensitive to scale
- bias correction (post-Lasso): use Lasso to get the "chosen" coefficients, then re-estimate coeffs by running them through OLS
 - note: post-inf model may show non-sig terms - ignore and keep
 - cannot fully trust inference of re-fitted model** (overfit)
- code: use `glmnet` to fit Ridge/Lasso and `cv.glmnet` for CV to find best λ (regularization strength)

```
cv_lambda <- cv.glmnet(
  x = X_train, y = y_train,
  alpha = 0/1, # 0 for Ridge, 1 for Lasso
  lambda = exp(seq(-5, 10, 0.1))) # range of lambda to try
lambda_min <- cv_lambda$lambda.min # lambda w/ min MSE
ridge_coef <- coef(cv_lambda, s = lambda_min) # get coeffs
ridge_min_predict <- predict(cv_lambda, newx = X_test,
                             s = lambda_min) # predict on new data
ridge_rmse <- rmse(preds = ridge_min_predict, actuals = Y_test)

# cv_lambda$lambda.1se give lambda bt 1 SE of min (sometimes better)
# note: use newx for reg-regression bc it requires matrix
```

- step-wise: is greedy (once var removed, not considered again) - only partial contribution considered, bad when num predictors $p > n$
- regularized methods: evaluate all variables jointly and removal not permanent, effective even when $n \ll p$

Cross-Validation

- hold out method: split into training and testing, create model on training, evaluate on testing set
 - limitation: estimate of test error rate depends on split (variable), tend to overestimate test error rate (fix with CV)
- k -fold CV: split data into k folds, leave 1 as validation set, train on $(k - 1)$ folds, estimate test MSE on validation test, let every fold be validation set once, report RMSE/MSE as average of each run
- Leave-one-Out CV (LOOCV): N folds with 1 obs each, train on $N - 1$ fold and test on left-out fold, do this N times, average (R)MSE at the end
 - resource intensive bc lots of training (good for small ds)
 - nearly unbiased param estimate but high variance

Logistic Regression

- used to: 1) estimate & test relations bt exp variables and **binary response**, 2) predict probability of a **binary** response (classifier)
- when response is binary (can only be 0 or 1), can't use SLR because it'll give predictions between 0 and 1, also outside of 0 and 1
 - binary response Y follows a Bernoulli distribution so $E(Y_i) = p_i$
- odds: prob positive over prob negative i.e odds = $p/(1 - p)$
- the model: after R, you'll get output for β_0, β_1, \dots , model is

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k \quad (1)$$

$$\frac{p}{1-p} = e^{\beta_0} \times e^{\beta_1 X_1} \times \dots \times e^{\beta_k X_k} \quad (2)$$

- p is the prob of response being 1, $\log(p/1-p)$ is called **log odds**
- (1) interpretation: for one unit increase in X_i (other constant), expect a β_i increase in **log odds** (similar interp to LR)
- (2) interpretation: if $e^{\beta_j} > 1$ - for one unit inc in X_j , the **odds** of the event occurring inc by factor of e^{β_j} (multiplicative)
 - note: when $e^{\beta_j} < 1$, it's easier to interp $1/e^{\beta_j}$ (flip target too)
 - ex. `exp.studentYes` = 0.524, then that means $1/0.524 = 1.908$, odds non-default of non-student is 90% more likely

```
log_reg <- glm(
  formula = y ~ x, data = data, family = binomial
) # formula like LR, family = binomial (logistic), or poisson
tidy(log_reg, exponentiate = FALSE) # = TRUE for exp version
```

- inference: same stuff as LR

$$H_0: \beta_j = 0 \quad H_a: \beta_j \neq 0 \quad z_j = \hat{\beta}_j / \text{SE}(\hat{\beta}_j) \sim N(0, 1)$$

- $CI = \hat{\beta}_j \pm z_{\alpha/2} \text{SE}(\hat{\beta}_j)$ (use `tidy(model, conf.int = TRUE)`)
- prediction: plug X into eq (either ver) and solve for probability p

```
log_odds <- predict(model, newdata = new_data, type = "link")
odds <- exp(log_odds) # exponentiate to get odds

# use type = "response" if you want straight probability
# overdispersion: IRL, data variance != model's assumed variance
# fix: estimate dispersion parameter phi to correct SE of estimators
# code: change family = "quasibinomial" (same coeffs, diff SE)
```

Poisson Regression

- used when response is count (i.e # of crabs) - assume $Y_i \sim \text{Pois}(\lambda_i)$
- the model: after R, you get

$$\log(\lambda) = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k \quad (3)$$

$$\lambda = e^{\beta_0} \times e^{\beta_1 X_1} \times \dots \times e^{\beta_k X_k} \quad (4)$$

- λ is the expected number of occurrences ($\log(\lambda)$ is log of that)
- (1) interpretation: inc in X_i by one unit (all else constant) results in a β_i increase in the log of expected counts
- (2) interpretation: if $e^{\beta_i} > 1$, for one unit increase in X_i , expected count of event by a factor of e^{β_i} (multiplicative)

Classification Metrics

- classification: log-reg give probability \hat{p}_i , we define a threshold p_0 where we'd predict 1 if $\hat{p}_i > p_0$, 0 otherwise (usually 0.5)
- error rate is your accuracy (# right / total)
 - training error rate likely underestimate out-of-sample error rate
 - use cross-validation to estimate out-of-sample error rate

```
cv_logistic <- cv.glm(glmfit = model, data = data,
                      K = 10, cost = self_defined_cost_function)
```

- confusion matrix: show you type of errors made by model
 - True Positive (TP): # obs correctly predicted as 1
 - False Positive (FP): # obs incorrectly predicted as 1 (truly 0)
 - True Negative (TN): # obs correctly predicted as 0
 - False Negative (FN): # obs incorrectly predicted as 0 (truly 1)
 - (pred, actual): TP = (1, 1), FP = (1, 0), FN = (0, 1), TN = (0, 0)
 - precision: PR = TP/(TP + FP)
 - accuracy: ACC = (TP + TN)/n

```
cm <- confusionMatrix(data=as.factor(model_preds), # col of 0/1
                      reference = as.factor(data$target), positive = "1")
```

- sensitivity and specificity: define relative measures
 - sensitivity: estimated prob of predicting 1 given true class is 1
 - math: SN = TP/TP + FN
 - crucial in scenarios where missing a positive case could have serious consequences (i.e medical diagnosis)
 - specificity: estimated prob of predicting 0 given true class is 0
 - math: SP = TN/TN + FP
 - crucial when it's important to be correct when identifying negatives (i.e drug screening - don't want to falsely accuse)
 - when dec threshold, SN inc and SP dec
- AUC and ROC: sometimes we're not sure of threshold p_0 to choose
 - evaluate pred performance of model for all val of $p_0 \in [0, 1]$
 - area under the curve (AUC) measures the classification ability of the model, ranges from 0 to 1, higher is better

```
ROC_output <- roc( # ROC TAKES PROBABILITY IN THE PREDICTOR
  response = data$target,
  predictor = predict(model, newdata, type = "response"))
plot(ROC_output, print.auc=TRUE) # actually plot
```

Regularized Logistic Regression

- can also use shrinkage methods in binary logistic regression
 - Regularized Loss = $\text{Loss}(\beta) + \lambda \|\beta\|_i$
 - $\|\cdot\|$ is norm, use $\|\cdot\|_1$ for Lasso, $\|\cdot\|_2$ for Ridge
- note: `glmnet` takes variables as matrices
 - object: formula of your model (-1 removes the intercept col)

```
X_train <- model.matrix(object = y ~ . - 1, data=training_split)
y_train <- as.matrix(training_split$target, ncol = 1)
# do the same for testing split to get X_test, y_test
```

- CV with regularized logisitic regression

```
cv_lambda_ridge <- cv.glmnet(x = X_train, y = Y_train,
                             alpha = 0, # Ridge regression
                             family = "binomial", # logistic regression
                             type.measure = "auc", nfold = 10)
ridge_max_AUC <- glmnet(x = X_test, y = Y_test, alpha = 0,
                       family = "binomial", lambda = lambda_min)
```

- get λ from `cv_obj$lambda.min/1se` which gives max AUC score
- sometimes, $\hat{\lambda}_{1SE}$ preferable bc we get much simpler model